



# PySpark Scenario-Based Interview Questions

## DAY 9 – Performance Optimization Basics



**Karthik Kondpak**  
**9989454737**

# PySpark Scenario-Based Interview

## Questions (Complete Notes Series)

### DAY 9 — Performance Optimization Basics (Partitions, Cache, Broadcast)

#### Concepts Covered Today

- What is partitioning and why it matters
- Repartition vs Coalesce
- Caching & Persistence
- Broadcast joins
- Real-world optimization mindset

#### Sample Scenario: Large Orders Dataset

Assume `orders_df` has **500 million records** and is used multiple times in the pipeline.

## Question 1: Check Number of Partitions

### ◆ Scenario

Before optimizing, you want to understand how data is distributed.

### PySpark Solution

```
orders_df.rdd.getNumPartitions()
```

### Explanation

- Too few partitions → under-utilized cluster
- Too many partitions → task scheduling overhead
- Default comes from `spark.sql.shuffle.partitions`

## Question 2: Repartition Data for Parallel Processing

### ◆ Scenario

Orders are heavily skewed. Repartition data to improve parallelism.



## PySpark Solution

```
optimized_df = orders_df.repartition(200,  
"customer_id")
```



## Explanation

- repartition() causes a **full shuffle**
- Used when increasing partitions or redistributing data



## Question 3: Reduce Partitions Using Coalesce

### ◆ Scenario

After heavy filtering, data size reduces drastically. Reduce partitions efficiently.



## PySpark Solution

```
final_df = optimized_df.coalesce(50)
```



## Explanation

- coalesce() avoids shuffle
- Best when **reducing** number of partitions



## Question 4: Cache Data Used Multiple Times

### ◆ Scenario

The same dataset is used for **multiple aggregations and joins**.



## PySpark Solution

```
orders_df.cache()  
orders_.df.count()#triggers cache
```



## Explanation

- Cache stores data in memory
- Avoids recomputation
- Use only when dataset is reused



## Question 5: Persist with Storage Level

### ◆ Scenario

Dataset is too large to fit fully in memory.



## PySpark Solution

```
from pyspark import StorageLevel  
  
orders_df.persist(StorageLevel.MEMORY_AND_DISK)
```



## Explanation

- Prevents OOM errors

- Frequently asked conceptual question

## Question 6: Optimize Join Using Broadcast

### ◆ Scenario

customers\_df is small (50K rows) and joined with massive orders\_df.

### PySpark Solution

```
from pyspark.sql.functions import broadcast

optimized_join_df = orders_df.join(
    broadcast(customers_df),
    "customer_id",
    "inner"
)
```



## Explanation

- Broadcast avoids shuffle
- One of the **most important interview answers**



**Let's build your Data  
Engineering journey  
together!**



Call us directly at: 9989454737



<https://seekhobigdata.com/>

