



PySpark Scenario-Based Interview Questions (Complete Notes Series)

DAY 14 – Structured Streaming



Karthik Kondpak
9989454737

PySpark Scenario-Based Interview Questions (Complete Notes Series)

DAY 14 — Structured Streaming (Real-Time Data Pipelines)

Concepts Covered Today

- What is Structured Streaming
- Micro-batch vs Continuous processing
- Sources & sinks (Kafka, Delta)
- Checkpointing & fault tolerance
- Event time, watermarking
- Streaming joins & aggregations

Scenario

You work for an Indian UPI / payments platform.

Stream: transactions_stream (from Kafka)

- transaction_id
- merchant_id
- amount
- event_time

Requirement: Calculate **real-time total amount per merchant every 5 minutes.**

What is Structured Streaming?

Definition

Structured Streaming treats streaming data as an **unbounded table** and processes it using **Spark SQL semantics**.

Key idea: “*Streaming = continuous incremental batch processing.*”

Question 1: Read from Kafka

PySpark Code

```
stream_df = spark.readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "broker:9092")
\
    .option("subscribe", "upi_transactions") \
    .load()
```

⌚ Question 2: Parse Streaming Data



Kafka value is JSON.
Scenario

```
from pyspark.sql.functions import from_json, col
from pyspark.sql.types import *

schema = StructType([
    StructField("transaction_id", StringType()),
    StructField("merchant_id", StringType()),
```

```
StructField("amount", DoubleType()),  
StructField("event_time", TimestampType())  
])  
  
parsed_df = stream_df.select(  
    from_json(col("value").cast("string"),  
schema).alias("data")  
).select("data.*")
```

Question 3: Windowed Aggregation

◆ Scenario

Calculate merchant-wise total every 5 minutes.

```
from pyspark.sql.functions import window, sum  
  
agg_df = parsed_df \  
.withWatermark("event_time", "10 minutes")\  
.groupBy(  
    window("event_time", "5 minutes"),  
    "merchant_id"
```

```
).agg(sum("amount").alias("total_amount"))
```

💧 Watermarking (Interview Favourite)

◆ Why Needed?

- Handles late arriving data
- Prevents unbounded state growth

Example: Allow events up to **10 minutes late**.

Question 4: Checkpointing (Fault Tolerance)

◆ Scenario

Ensure exactly-once processing.

```
query= agg_df.writeStream \  
    .format("delta") \  
    .outputMode("append") \  
    .option("checkpointLocation", "/chk/upi_txn") \  
    .start("/delta/merchant_agg")
```

Output Modes

Mode	Use Case
append	New rows only
update	Changed rows
complete	Full result table

↻ Question 5: Streaming Join (Advanced)

◆ Scenario

Join transaction stream with static merchant data.

```
joined_df = parsed_df.join(  
    broadcast(merchants_df),  
    "merchant_id"  
)
```



**Let's build your Data
Engineering journey
together!**

 Call us directly at: 9989454737

 <https://seekhobigdata.com/>

