



PySpark Scenario-Based Interview Questions (Complete Notes Series)

**DAY 22 – Spark Memory
Management & GC Tuning**



Karthik Kondpak
9989454737

PySpark Scenario-Based Interview

Questions (Complete Notes Series)

DAY 22 — Spark Memory Management & GC Tuning

Concepts Covered Today

- Spark memory architecture
- Execution vs Storage memory
- On-heap vs Off-heap memory
- Garbage Collection (GC) basics
- GC tuning strategies Real Indian
- production incidents

Spark Memory Architecture (MUST KNOW)

Executor JVM Memory



 └ Execution Memory

```
|   └ Shuffles  
|   └ Joins  
|   └ Aggregations  
|   └ Sorts  
|  
|  
└ Storage Memory  
|   └ Cached DataFrames  
|   └ Broadcast variables  
|   └ Persisted RDDs  
|  
└ Reserved Memory (System)
```

Execution & Storage memory **share space dynamically.**

Key Spark Memory Configurations

```
spark.executor.memory  
spark.executor.memoryOverhead  
spark.memory.fraction          (default 0.6)  
spark.memory.storageFraction (default 0.5)
```

Scenario

A Spark job processing **IRCTC ticket booking data**:

- Heavy joins
- Cached lookup tables
- Frequent executor crashes

Error seen:

`ExecutorLostFailure (OOM)`

Why Spark Jobs Fail with OOM?

- Large shuffles
- Uncontrolled caching
- Broadcasting big tables
- Skewed partitions
- Poor GC behavior

Execution vs Storage Memory (COMMON QUESTION)

Area	Used For	Priority
Execution	Joins, shuffles	High
Storage	Cache, persist	Lower

Execution memory can evict cached data.

Caching & Persistence Best Practices

`df.persist(StorageLevel.MEMORY_AND_DISK)`

Cache only reused data

Do not cache large intermediate results

Garbage Collection (GC) in Spark

- ◆ What is GC?

GC cleans unused JVM objects to free memory.

Problem:

- Excessive GC pauses slow jobs dramatically

Signs of GC Issues

- Tasks slow despite low CPU usage
- Executors alive but not progressing
- Long GC time in Spark UI

GC Tuning

- ◆ **Use G1GC (Recommended)**

```
spark.executor.extraJavaOptions = -XX:+UseG1GC
```

- ◆ **Tune Heap Occupancy**

```
-XX:InitiatingHeapOccupancyPercent=35
```

On-Heap vs Off-Heap Memory

Type	Pros	Cons
On-Heap	Managed by GC	GC overhead
Off-Heap	Less GC	Manual tuning

`spark.memory.offHeap.enabled=true`

`spark.memory.offHeap.size=4g`



**Let's build your Data
Engineering journey
together!**



Call us directly at: 9989454737



<https://seekhobigdata.com/>

