

Spark Optimization Topic

**Bucketing – How It Reduces
Shuffles and Speeds Up Joins**



Karthik Kondpal

Seekho Bigdata Institute www.seekhobigdata.com 9989454737

9989454737

Day 13 — Spark Optimization Topic

🔥 Bucketing — How It Reduces Shuffles and Speeds Up Joins

Bucketing is one of Spark's most **powerful shuffle-reduction techniques**, but also one of the **least understood**.

If you struggle with large JOINs, slow aggregations, or skewed data

—

Bucketing can save your job.

What Is Bucketing?

Bucketing **divides data into fixed number of buckets** using a **hash function** on a column.

Example:

```
BUCKET 4 ON customer_id
```

means Spark will hash the customer_id values → assign rows into 4 fixed buckets.

Unlike partitioning:

- ✓ Buckets are **pre-distributed**
- ✓ Every bucket is same size (hash-based)
- ✓ Useful for JOIN and GROUP BY
- ✓ Works across multiple tables

📌 Why Bucketing Matters

- ✓ Joins become faster

If two tables are bucketed on the same key, Spark **avoids shuffles**.

- ✓ GroupBy becomes faster

Spark can skip repartitioning and reuse existing buckets.

- ✓ Reduces memory pressure

Less shuffle → fewer shuffle partitions → fewer spills.

✓ Works on massive datasets

Used heavily in:

- E-commerce (Flipkart, Amazon India)
- Telecom
- Banking
- Ad-tech

Bucketing Example (Syntax)

Creating a bucketed table

```
df.write \
    .bucketBy(8, "customer_id") \
    .sortBy("customer_id") \
    .format("parquet") \
    .saveAsTable("bucketed_customers")
```

Reading it

```
df = spark.table("bucketed_customers")
```

Why SORT BY?

Sorting inside the bucket:

- ✓ speeds up merges
- ✓ reduces downstream sort work
- ✓ helps in window functions, joins, and groupBy

Scenario

You work at Flipkart India, handling customer transactions.

Two tables:

1. customers (100M rows)
2. orders (1.2B rows)

You frequently join:

```
SELECT *
FROM customers c
JOIN orders o
```

```
ON c.customer_id = o.customer_id
```

If both tables are bucketed on customer_id with same number of buckets,

Spark:

- ✓ Avoids shuffle
- ✓ Reads matching buckets directly

Pipeline runs 3–5× faster

- ✓ Reduces cluster cost
- ✓

How Bucketing Removes Shuffle

Normal join physical plan:

Exchange (HashPartitioning)

Sort Join

With bucketing:

Scan (Bucketed)

SortMergeJoin

Notice:

No Exchange → No shuffle

- ✓ Already bucketed → Already partitioned
- ✓ Spark just sorts inside bucket → Joins → Done

When to Use Bucketing

Use bucketing when:

- ✓ You join the **same two big tables repeatedly**
- ✓ Same join key is used often (customer_id, product_id, mobile_no)
- ✓ Data is stable (not updated very frequently)
- ✓ You want predictable distribution

When NOT to Use Bucketing

Bucketing is NOT ideal if:

- ✗ Data is highly dynamic
- ✗ You don't manage tables (external sources)
- ✗ You don't know join keys ahead of time
- ✗ You rarely join those tables



**Let's build your Data
Engineering journey
together!**



 Call us directly at: 9989454737



<https://seekhobigdata.com/> Seekho Bigdata Institute www.seekhobigdata.com 9989454737