

Spark Optimization Topic



Karthik Kondpak

Seekho Bigdata Institute www.seekhobigdata.com 9989454737
9989454737

Day 14 — Spark Optimization Topic

Joins — Choosing the Right Join for Speed (with Physical Plan Explained)

In PySpark, JOINS are the **most expensive operations** because they often trigger:

- ✓ Shuffle

- ✓ Sort

- ✓ Network transfer

- ✓ Data skew

- ✓ High memory usage

Choosing the **right join strategy** can reduce runtime from **minutes** → **seconds**.

6 Join Types Every Data Engineer Must Know

Spark uses different physical join strategies depending on data size and configuration.

Broadcast Hash Join (BHJ)

Fastest join when one table is small.

✓ When to use

When one table is **≤ 10 MB** (default broadcast threshold).

```
df1.join(broadcast(df2), "id")
```

✓ Physical Plan

BroadcastHashJoin

```
[ *BroadcastExchange* ]
```

✓ Why it's fast

- No shuffle
- Small table copied to executors
- Big table scanned locally

Shuffle Hash Join (SHJ)

Used when both tables are medium-sized and hashable.

✓ Physical Plan

HashExchange

HashExchange

ShuffledHashJoin

✓ When to use

- Tables fit in memory
- Join key is well distributed
- No skew

⚠ If skew exists: performance drops sharply.

Sort Merge Join (SMJ)

Default join for large datasets.

✓ When to use

- Both tables huge
- Join key cannot be broadcast
- Sorting and merging required

✓ Physical Plan

Exchange (HashPartitioning)

Sort

Exchange (HashPartitioning)

Sort

SortMergeJoin

✓ Why expensive

- Heavy shuffle
- Heavy sort
- Large data movement

Cartesian Join (CROSS JOIN)

Produces $N \times M$ rows — extremely expensive.

Use only when **logically required**.

✓ **Physical Plan**

CartesianProduct

⚠ Avoid at all cost unless business logic demands it.

Skew Join Optimization (AQE)

When data is skewed:

- One key has too many rows
- Executors take long
- Entire job slows

AQE fixes it by:

✓ Splitting skewed partitions

✓ Broadcasting smaller side

✓ Using hybrid join strategy

✓ **Physical Plan**

AdaptiveSparkPlan

SplitSkewedPartition

BroadcastHashJoin

Bucketed Sort Merge Join (Ultra-fast for large data)

If two tables are **bucketed on the same key AND same bucket count**, Spark avoids shuffle.

✓ **Physical Plan**

Scan (Bucketed)

SortMergeJoin

✓ No Exchange

✓ No repartitioning

✓ Much faster

Join Selection Matrix — Which Join Should You Use?

Table Sizes	Best Join	Why
Small + Large	Broadcast Hash Join	No shuffle
Medium + Medium	Shuffle Hash Join	Faster than SMJ
Large + Large	Sort Merge Join	Handles big data
Pre-bucketed tables	Bucketed SMJ	Zero shuffle
Skewed keys	AQE + Skew Join	Splits heavy partitions

Physical Plan Comparison

Broadcast Join

`BroadcastHashJoin`

`+ - BroadcastExchange`

+-- Project / Filter / Scan

Sort Merge Join

SortMergeJoin

+-- Sort

+-- Exchange

+-- Scan

Shuffle Hash Join

ShuffledHashJoin

+-- Exchange

+-- Exchange

Scenario Example :

You work at **Swiggy India**, processing daily restaurant orders.

Tables:

- **orders** → 900M rows
- **restaurants** → 40 MB

Best join strategy: Broadcast Join

```
df_orders.join(broadcast(df_restaurants),  
"restaurant_id")
```

✓ Saves shuffle

✓ Saves sort

Saves network cost

✓ Ideal for daily pipelines

✓

Expert Tip

Before joining ALWAYS check size:

```
df.count()  
spark.conf.get("spark.sql.autoBroadcastJoinThreshold")
```

And check for skew:

```
df.groupBy("join_key").count().orderBy(desc("count"))
```



**Let's build your Data
Engineering journey
together!**



Call us directly at: 9989454737



<https://seekhobigdata.com/>

