

# File Size Optimization



**Karthik Kondpak**

Seekho Bigdata Institute [www.seekhobigdata.com](http://www.seekhobigdata.com) 9880454737

**9989454737**

# Day 15 — Spark Optimization Topic

## 🔥 File Size Optimization — How Small

### Files Kill Spark Performance

Small files are one of the **biggest silent killers** in Spark jobs.

They slow down queries, increase cost, and overload the cluster without you even noticing.

If your data lake has **thousands of small Parquet/Delta files**, Spark will struggle — even if the data size is small.

### Why Small Files Hurt Performance

Small files create:

**Too many tasks → Too much scheduling overhead**

Spark creates **1 task per file** per partition.

Example:

10,000 Parquet files = 10,000 tasks.

Most tasks finish in milliseconds, but **scheduler overhead** becomes huge.

### **Large Shuffle and Skew**

Small files → small partitions → imbalance in size.

Some partitions = 5 KB

Some partitions = 500 MB

This causes:

- Slow executors
- Long-running tasks
- Skew during shuffle

### **High Metadata Load on Hive Metastore / Glue Catalog**

Each file entry = one metadata entry.

100K small files → Metadata queries become slow.

### **Slow reads & writes**

Spark reads multiple tiny row groups → more I/O operations.

## Delta Lake Transaction Log growth

Each write adds one entry.

Small files → Too many commits → Slow time travel & vacuum.

## Ideal File Size for Spark

Format	Ideal File Size
Parquet	256 MB – 1 GB
ORC	256 MB – 512 MB
Delta	256 MB – 1 GB

Anything below **16 MB** is considered a "small file" problem.

## Scenario :

You work at **Zomato India**.

A pipeline ingests **600 MB** of daily orders, but writes **3,000 Parquet files of 200 KB each.**

Later when analysts query this table:

```
select * from orders where city = 'Mumbai'
```

Spark must:

- Open 3,000 files
- Read metadata per file
- Create thousands of tasks

The query becomes 10× slower.

## How Spark Creates Small Files

You get small files when:

- Writing without partitioning correctly
  - Writing with too many partitions
  - Upstream system produces micro-batches
  - Using **foreachBatch** or **streaming jobs**

- Using **overwrite** in loops
- Using **repartition(N)** incorrectly

## How to Fix Small Files

### Repartition before writing

Best method.

```
df.repartition(50).write.format("delta").save("/path")
```

Choose partition count using:

`total_data_size / 256MB`

### Use **coalesce()** for reducing partitions

Faster than repartition (no shuffle).

```
df.coalesce(20).write.parquet("/path")
```

## Use OPTIMIZE (Delta Lake Only)

```
OPTIMIZE orders  
WHERE order_date >= '2024-11-01'
```

This compacts small files → Large healthy files.

## Use Auto-optimize + Auto-compaction (Databricks)

```
spark.conf.set("spark.databricks.ckss.enabled", "true")  
spark.conf.set("spark.databricks.ckss.optimizeWrite.enabled", "true")
```

## **Use `maxRecordsPerFile` to avoid extremely large files**

```
df.write.option("maxRecordsPerFile",  
1000000).parquet("/path")
```

## **Compact files manually (Apache Spark)**

(

```
spark.read.format("parquet").load("/path")  
    .repartition(20)  
    .write.mode("overwrite")  
    .parquet("/path_temp")
```

)

Then swap folders.

## Before & After Example

Condition	Spark Tasks	Runtime
Before (3000 small files)	3000+ tasks	20 mins
After compaction (20 files)	20 tasks	2 mins

10× FASTER — same data, better file layout.



**Let's build your Data  
Engineering journey  
together!**



Call us directly at: 9989454737



<https://seekhobigdata.com/>

Seekho Bigdata Institute www.seekhobigdata.com 9989454737

