



# 50+ PySpark Interview Questions — With Answers



**Karthik Kondpak**  
**9989454737**

# 50+ PySpark Interview Questions — With Answers



## 1. What is PySpark? Explain in detail.

### Definition

PySpark is the **Python API for Apache Spark**, which allows developers to write Spark applications using Python instead of Scala or Java.

PySpark interacts with the underlying Spark engine through:

- **Py4J (Python → JVM communication bridge)**
- **Apache Arrow (for optimized columnar data transfer, especially in Pandas UDF)**

### Why PySpark?

- Easy to code for Python developers
- Integrates well with ML libraries (Pandas, NumPy, MLlib)
- Supports distributed processing on huge datasets

- Optimized using Catalyst + Tungsten

## Internal Flow

Python Code → Py4J → JVM → Spark Execution Engine → Executors

### Real-Time Use Case

Processing 2TB clickstream logs from Flipkart across 300 executors in parallel.

## ✓ 2. What is SparkSession? Why is it important?

### Definition

SparkSession is the **entry point** to all Spark functionalities.

Before Spark 2.0 → We had:

- SQLContext
- HiveContext
- SparkContext

After Spark 2.0 → All merged into **SparkSession**.

## **What does SparkSession manage?**

- SQL Query Engine
- DataFrame creation
- Catalyst Optimizer
- Configuration
- Catalog (database/tables metadata)
- Communication with cluster manager

## **Example**

```
spark = SparkSession.builder \
    .appName("MyApp") \
    .config("spark.sql.shuffle.partitions", 100)
\
    .getOrCreate()
```

## **Why is it important?**

Without SparkSession, no DataFrame operations can run.



### 3. Explain Spark Architecture in detail.

#### Components

##### 1. Driver Program

- a. Creates SparkSession
- b. Converts code into DAG
- c. Schedules tasks

##### 2. Cluster Manager

- a. YARN / Kubernetes / Standalone

##### 3. Executors

- a. JVM processes running tasks
- b. Store data in memory
- c. Send results back to driver

#### Flow

Python Code → Driver → DAG Scheduler → Task Scheduler →  
Executors → Results → Driver

#### Example

If dataset = 1TB, and 100 executors are available →

Each executor processes **10GB** in parallel.



## 4. Explain DAG (Directed Acyclic Graph) in Spark.

### Definition

DAG = Representation of your transformations as a flow of operations with no cycles.

Example DAG:

Read → Filter → Select → GroupBy → Write

### Stages

A wide transformation causes a **new stage**:

Stage 1 → narrow ops

Stage 2 → shuffle ops

Stage 3 → write ops

### Why important?

- Allows optimization

- Parallel execution
- Fault tolerance (can recompute lineage)

## 5. What are Transformations and

### Actions? Give deep explanation.

#### Transformations (Lazy)

- Do NOT execute immediately
- Build logical plan
- Are **lazy evaluated**

Examples:

- map, filter, select
- join
- groupBy
- withColumn

#### Actions (Trigger execution)

Examples:

- show

- count
- collect
- write

### Why laziness?

- Spark optimizes full DAG end-to-end
- Reduces computation
- Avoids unnecessary processing

## ✓ 6. Explain Narrow vs Wide Transformations.

### Narrow (Fast)

- Parent partition → single child partition
- No shuffle
- Executed within same stage

Examples:

- map
- filter

- withColumn
- select

### **Wide (Slow)**

- Requires shuffle
- New stage is created

Examples:

- groupBy
- join
- distinct
- orderBy

### **Why wide transformations are expensive?**

- Data movement between nodes
- Write to disk
- Network IO
- Serialization/deserialization

## 7. What is Catalyst Optimizer? Explain in depth.

Catalyst Optimizer is Spark's **rule-based + cost-based SQL optimizer**.

### Optimization Phases

#### 1. Analysis

- a. Resolve column names
- b. Type checking

#### 2. Logical Optimization

- a. Filter pushdown
- b. Constant folding
- c. Predicate reordering
- d. Null propagation
- e. Projection pruning

#### 3. Physical Planning

- a. Choose best join strategy
- b. Sort-merge join vs hash join

#### 4. Code Generation

- a. Whole stage code generation

## Why Catalyst is powerful?

- Minimizes CPU operations
- Minimizes IO
- Uses statistics (cost-based optimizer)

## ✓ 8. What is Tungsten Engine? Explain technically.

Tungsten introduces **low-level optimizations** for speed:

### Main Enhancements

1. Off-heap memory
2. Binary memory format
3. Whole-Stage Code Generation
4. Eliminate JVM object overhead
5. Columnar caching

### Result

- Faster joins
- Faster aggregations

- Less GC overhead

## 9. Explain Whole-Stage Code

### Generation with example.

Spark combines multiple execution operators into a **single optimized function**.

Without WSCG:

Row → Filter → Map → Reduce (multiple function calls)

With WSCG:

Single Java code block

#### Benefit

- 3–5x faster execution
- Fewer CPU cycles
- Better pipeline optimization



## 10. Explain Broadcast Join in detail.

### Definition

Broadcast join sends a small table to **every executor node**.

### When to use?

- Table size < 100–500 MB
- Dimension/LUT tables
- When avoiding shuffle is key

### Example

```
from pyspark.sql.functions import broadcast
df1.join(broadcast(df2), "id")
```

### Internal Working

- Table df2 is serialized
- Copied to each executor
- Join becomes local operation → no shuffle

## 11. What is Data Skew? How to detect and solve?

### Definition

Uneven distribution of keys during join/groupBy.

Example:

customer\_id = 99999 → 50 million records

customer\_id = others → only 10,000 records

Executor handling 99999 becomes slow → job fails.

### Detection

1. Spark UI → skewed partition sizes
2. High task duration in one executor

### Fix

#### 1. Salting

```
df = df.withColumn("salt", (rand()*10).cast("int"))
```

#### 2. Broadcast smaller table

### 3. Skew join hints

```
df1.join(df2.hint("skew"), "id")
```

4. Increase partitions

5. Range partitioning



## 12. What is Partitioning in Spark?

### Explain deeply.

#### Definition

The way Spark splits data into parallel chunks.

#### Why important?

- Determines parallelism
- Impacts shuffle
- Affects memory distribution
- Impacts file writing

#### Types

##### 1. Hash partitioning

- 2. Range partitioning**
- 3. Custom partitioning**

#### Example

```
df.repartition(200, "cust_id")
```

### ✓ 13. Difference between repartition() and coalesce().

Feature	repartition()	coalesce()
Shuffle	Yes	No
Use Case	Increase or decrease partitions	Decrease only
Performance	Expensive	Cheap
Guarantees	Even distribution	Uneven distribution possible

### ✓ 14. Explain Window Functions in PySpark with examples.

Used for:

- Ranking

- Moving averages
- Running totals
- LAG/LEAD
- Time-based aggregations

### **Example**

```
from pyspark.sql.window import Window
from pyspark.sql.functions import lag

w = Window.partitionBy("dept").orderBy("salary")

df.withColumn("prev_salary", lag("salary",
1).over(w)).show()
```

### **Real-time Use Case**

Finding last 3 transactions per customer.

## ✓ 15. Explain explode(), poseplode(), explode\_outer().

**explode**

Flatten array/map column → emits one row per element.

**poseplode**

Emits **position + element**.

**explode\_outer**

Same as explode but preserves null entries.

**Example**

```
df.withColumn("item", explode("items"))
```

Perfect — I will now continue with **Q16–Q50**, but with **EVEN MORE DETAIL**, suitable for **creating full PySpark notes / eBook / interview guide**.

These answers will include:

✓ Internal architecture

✓ Real-time data engineering scenarios

Best practices

✓ PySpark code

✓ Hidden concepts interviewers expect

✓ Spark UI + performance notes

✓

## ✓ 16. How do you read Nested JSON in PySpark? Explain in depth.

### ★ Why nested JSON is difficult?

- JSON contains complex structures (arrays, structs, maps, nested structs).
- Schema is hierarchical.
- Flattening JSON creates wide tables with many nested levels.

## ★ How Spark handles JSON

Spark's JSON reader:

- Infers schema (if enabled)
- Parses hierarchical fields into `StructType`, `ArrayType`, `MapType`
- Allows exploding & flattening

## ★ Example JSON

```
{  
    "id": 101,  
    "name": "Shivani",  
    "orders": [  
        {"order_id": 1, "amount": 2000},  
        {"order_id": 2, "amount": 3500}  
    ],  
    "address": {  
        "city": "Mumbai",  
        "pin": 400001  
    }  
}
```

## ★ Read JSON

```
df = spark.read.json("/path/nested.json",
multiLine=True)
df.printSchema()
```

## ★ Flatten JSON

```
from pyspark.sql.functions import col, explode
df2=df.withColumn("order", explode("orders")) \
.select(
    "id",
    "name",
    col("address.city").alias("city"),
    col("address.pin").alias("pincode"),
    col("order.order_id"),
    col("order.amount")
)
```

## ★ Use Case (e-commerce)

Flipkart API logs, Shopify webhooks, Razorpay payment events.

## ★ Optimization Tips

- Disable inferSchema and provide schema manually
- Always use multiLine=True for multi-row JSON objects

## ✓ 17. What are UDFs? Why are they slow?

### Detailed explanation.

#### ★ What is a UDF?

A User Defined Function lets you write Python functions and apply them to Spark DataFrames.

```
from pyspark.sql.functions import udf
```

```
def multiplier(x): return x * 2
udf_mult = udf(multiplier)
df.withColumn("result", udf_mult("value"))
```

#### ★ Why are Python UDFs slow?

Reason	Explanation
<b>Breaks Catalyst Optimization</b>	Spark cannot optimize custom Python logic
<b>Row-by-row execution</b>	UDF processes each row individually

<b>Serialization cost</b>	Python → JVM → Python
<b>Cannot use vectorization</b>	No batch operations
<b>Not memory-efficient</b>	Creates many Python objects

### ★ Better Alternatives

- **Built-in functions**
- **SQL expressions**
- **Pandas UDFs**
- **Scala UDF for speed-critical tasks**

### ★ Real-Time Example

Banking company applying regex-based fraud detection logic on 600M records — Python UDF was 20x slower, replaced with **Spark SQL built-ins**.

## ✓ 18. What are Pandas UDFs? Why are they faster?

### ★ Definition

Pandas UDFs (also known as vectorized UDFs) use **Apache Arrow** for **columnar, batch-based** execution.

## ★ Why faster?

Feature	Python UDF	Pandas UDF
Execution	Row-by-row	Batch-vectorized
Data transfer	Python $\leftrightarrow$ JVM (slow)	Arrow-optimized columnar
Speed	Slow	5x–20x faster
Optimized by Catalyst?	No	Yes

## ★ Example

```
import pandas as pd
from pyspark.sql.functions import pandas_udf
from pyspark.sql.types import DoubleType

@pandas_udf("double")
def multiply(col: pd.Series) -> pd.Series:
    return col * 10

df.withColumn("new_col", multiply("value"))
```

## ★ When to use?

- ML feature engineering
- Complex math operations
- Statistical calculations

## ✓ 19. What is cache() vs persist()? Deep explanation.

### ★ cache()

Default storage = MEMORY\_ONLY  
df.cache()

### ★ persist()

Allows custom storage level:

```
from pyspark import StorageLevel  
df.persist(StorageLevel.MEMORY_AND_DISK)
```

### ★ Storage options

- MEMORY\_ONLY
- MEMORY\_AND\_DISK
- MEMORY\_ONLY\_SER
- MEMORY\_AND\_DISK\_SER
- DISK\_ONLY

### ★ When to use caching?

- ✓ Data reused multiple times
- ✓ Iterative algorithms
- ✓ Avoid recomputation of expensive transformations

### ★ When NOT to use caching?

- ✗ One-time use DataFrames
- ✗ Cache too many large DataFrames → executor OOM

### ★ Real example

Reusing customer dataset 8 times during product recommendation creation → caching reduces execution from **50 min** → **7 min.**

## ✓ 20. Explain the top 10 PySpark

### Optimization Techniques (very detailed).

#### ★ 1. Avoid Wide Transformations

groupBy, join, distinct → cause shuffle

Shuffle = disk IO + network IO → slowest operation in Spark.

## ★ 2. Use Broadcast Join

```
df1.join(broadcast(df2), "id")
```

Avoids shuffle.

## ★ 3. Reduce Shuffle Partitions

Default partitions = 200 / 300 (too high)

Set:

```
spark.conf.set("spark.sql.shuffle.partitions", 50)
```

## ★ 4. Use Filter Early (Predicate Pushdown)

Filter before join/groupBy:

```
df.filter("year = 2024")
```

## ★ 5. Avoid Python UDFs

Use built-in or Pandas UDFs.

## ★ 6. File Format Choice

Use:

✓ Parquet

✓ ORC

✓ Delta

Avoid CSV/JSON during processing.

## ★ 7. Avoid Collect()

collect() pulls entire dataset to driver → OOM.

## ★ 8. Repartition Based on Join Key

```
df1.repartition("id").join(df2.repartition("id"))
```

## ★ 9. Use checkpoint() to break lineage

Very long DAG may cause stack overflow:

```
df.checkpoint()
```

## ★ 10. Use AQE (Adaptive Query Execution)

Automatically adjusts:

- join strategy
- shuffle partitions
- skew optimization

```
spark.conf.set("spark.sql.adaptive.enabled", True)
```

## 21. Explain Predicate Pushdown in Spark.

### What is Predicate Pushdown?

Spark pushes filters to the file reader **before** reading full data.

Example:

```
SELECT * FROM sales WHERE year = 2024
```

Spark reads **only partitions where year=2024**, not entire dataset.

### Works with

- Parquet
- ORC
- Delta
- JDBC (limited)

### Does NOT work with

- CSV
- JSON

### **Code Example**

```
df.filter("year = 2023")
```

### **Benefit**

- Reduces IO
- Reduces scan time
- Faster joins/aggregations

## **22. Explain Column Pruning.**

Spark reads only required columns.

### **Example**

```
df.select("id")
```

If file has 100 columns, Spark reads **1 column** only.

### **Benefit**

- Faster scan
- Less memory usage

- Small memory footprint

## ✓ 23. What is Delta Lake? Explain in detail.

★ Delta Lake = Parquet + Transaction Log (.delta)

★ Features

Feature	Description
ACID transactions	Safe writes in big data
Time travel	Query old versions
Schema evolution	Auto add columns
Schema enforcement	Prevents corrupted writes
Upsert/Merge	SQL MERGE support
Faster reads	File pruning + metadata optimization

★ Example

```
df.write.format("delta").save("/path")
df = spark.read.format("delta").load("/path")
```

★ Real-Time Use Case

Data Lakehouse architecture on Databricks / AWS EMR.

## 24. Explain Delta Merge (UPSERT) with real-time example.

### Scenario

Updating daily incremental data to customer master table.

### Code

```
from delta.tables import DeltaTable

dt = DeltaTable.forPath(spark, "/delta/customer")

(
    dt.alias("t")
    .merge(
        updates.alias("u"),
        "t.id = u.id"
    )
    .whenMatchedUpdateAll()
    .whenNotMatchedInsertAll()
    .execute()
```

)

## ★ Optimizations

- Z-Ordering
- Vacuum
- OPTIMIZE command

## ✓ 25. What is Schema Evolution vs Schema Enforcement?

### Schema Evolution (Allow changes)

Spark automatically updates schema.

```
spark.conf.set("spark.defaultTableCheck.enabled.autoMerge",  
    true)
```

### ★ Schema Enforcement (Reject mismatches)

Spark blocks bad schemas.

Example error:

```
column price expected INT got STRING
```

## ★ Use Cases

- Changing JSON structures
- Adding new fields in Kafka streams
- Modifying CDC tables

## ✓ 26. How to optimize join performance?

### (Very detailed)

#### ★ Best Practices

**Broadcast small table**

**Repartition both tables on join key**

```
df1.repartition("id").join(df2.repartition("id"),  
"id")
```

**Use Sort-Merge Join only when sorted**

**Avoid Skewed Keys**

**Use Bucketing for large static tables**

**Use AQE → auto tune joins**



## 27. What is File Pruning?

Spark reads only **required partitions** based on partition column.

Directory structure:

/year=2024/month=07/day=01

Query:

```
WHERE year = 2024 AND month = 7
```

Spark does NOT read all folders → only required ones.



## 28. What are Accumulators?

### ★ Definition

Write-only variables used in:

- Counting
- Debugging
- Monitoring job metrics

### ★ Example

```
acc =  
spark.sparkContext.longAccumulator("error_count")
```

### ★ Limitation

Executors can only increment, cannot read.

## ✓ 29. What are Broadcast Variables?

### ★ Definition

Read-only variable distributed to executors.

Used for:

- Lookup tables
- Config constants

### Example

```
bv = spark.sparkContext.broadcast({"IN": "India",  
"US": "United States"})
```

## 30. How to handle large datasets efficiently in PySpark?

- ✓ Use Parquet/ORC/Delta
- ✓ Avoid collect()
  - Write partitioned data
- ✓ Cache wisely
- ✓ Avoid Python UDFs
- ✓ Push filters early
- ✓ Use correct cluster config
- ✓ Use coalesce while writing small files
- ✓

✓ Avoid small file problem

✓ Use efficient serialization (Kryo)

Great — continuing with **Q31–Q50**, in the same ultra-detailed, note-style format ready for conversion into notes/ebook/carousel slides. Each question includes: concept, internals, code snippets, real-world use cases, performance tips, and interview-ready talking points.

## ✓ **31. What is Structured Streaming?**

**Explain internals, modes, and code.**

**Concept**

Structured Streaming is Spark's high-level API for stream processing built on top of the DataFrame/Dataset APIs. It treats streaming data as an **unbounded table** and provides exactly-once semantics for many sources/sinks.

## Modes

- **Micro-batch** (default): processes data in small batches (e.g., every 1s).
- **Continuous** (experimental): low-latency processing (ms), limited operators.

## Internals

- Stream is represented as a logical plan; continuous mini DAGs executed periodically.
- Checkpointing stores progress & state.
- Watermarking handles late data and state eviction.
- State store (RocksDB/Memory) manages stateful operators.

## Example (micro-batch)

```
df = spark.readStream.format("kafka") \
    .option("kafka.bootstrap.servers","host:9092")
    \
    .option("subscribe","topic").load()

parsed = df.selectExpr("CAST(value AS STRING) as
json") \
    .select(from_json(col("json"),
```

```

schema).alias("data")) \
    .select("data.*")

agg = parsed.groupBy(window("event_time", "10
minutes"), "country").count()

query = agg.writeStream \
    .outputMode("append") \
    .format("parquet") \
    .option("path", "/mnt/stream_output") \
    .option("checkpointLocation",
"/mnt/checkpoints/agg") \
    .start()

```

## Use cases

Real-time analytics, fraud detection, live dashboards, CDC ingestion.

## Performance tips

- Tune trigger interval and processing time.
- Use watermarking to drop old state.
- Keep state small; use mapGroupsWithState cautiously.

- Persist state store to durable storage.

## 32. What is Watermarking and why use it?

### Concept

Watermarking tells Spark how long to wait for late data for event-time aggregations. It bounds state size and allows emitting results while tolerating late arrivals up to the watermark threshold.

### How it works

`withWatermark(eventTimeCol, "delay")` marks the maximum lateness; Spark drops state older than `maxEventTime-delay` .

### Example

```
agg = events.withWatermark("event_time", "10
minutes") \
    .groupBy(window("event_time", "5
minutes"), "user_id") \
    .count()
```

## Use cases

Session windows, rolling aggregates, page-view counts with late logs.

## Pitfalls

- If watermark too small → drop valid late events.
- Too large → state grows and memory pressure rises.

## 33. What is checkpointing (streaming & batch) and why is it needed?

### Purpose

- **Streaming:** Save progress (offsets, state, metadata) to recover from failures and provide exactly-once guarantees.
- **Batch:** Shorten lineage (checkpointing RDDs/DataFrames) to avoid very long DAGs and stack overflow, or to make stable points for iterative algorithms.

### **Example (streaming)**

```
query = df.writeStream \
    .format("parquet") \
    .option("checkpointLocation",
"/mnt/checkpoint/stream1") \
    .start()
```

### **Example (batch)**

```
spark.sql("CREATE TABLE IF NOT EXISTS checkpoint_intDir AS SELECT * FROM parquet.`/tmp/checkpoint`")
df.checkpoint()
```

### **Tips**

- Always provide checkpointLocation for production streams.
- Use durable storage (S3/ADLS/DBFS) for checkpoint directories.
- Periodically clean up old checkpoints carefully when upgrading logic.

## 34. Explain Adaptive Query Execution (AQE) and how it helps.

### Concept

AQE allows Spark to modify the physical execution plan at runtime using statistics collected during execution. It can:

- Dynamically switch join strategies.
- Coalesce shuffle partitions based on actual partition sizes.
- Handle skewed data by splitting/handling large partitions.

### How to enable

```
spark.conf.set("spark.sql.adaptive.enabled", "true")
```

### Benefits

- Better resource utilization.
- Avoids over/under partitioning.
- Automatic skew mitigation.

### Caveats

- Slight overhead for statistics collection.

- Works best with Parquet/ORC/Delta with file size stats.

## ✓ 35. How to handle data skew —

### techniques with code & pros/cons

#### Detection

- Spark UI: one or few tasks take much longer.
- Highly imbalanced partition sizes.

#### Techniques

##### 1. Salting

- a. Add random suffix to skewed keys then join/group, then aggregate/sum back.

```
from pyspark.sql.functions import rand, concat, lit
df_saltd = df.withColumn("salt",
(rand()*10).cast("int"))
df_saltd = df_saltd.withColumn("salted_key",
concat(col("key"), lit("_"), col("salt")))
```

- b. Pros: simple. Cons: increases data size & complexity.

## **2. Broadcast the smaller side**

- a. If one table is small.

```
df_large.join(broadcast(df_small), "key")
```

- b. Pros: avoids shuffle. Cons: limited by broadcast size.

## **3. Skew join hint / AQE**

- a. Let Spark auto-handle when enabled or use hints.
- b. Pros: automated. Cons: may not always pick perfect strategy.

## **4. Range partitioning**

- a. Pre-partition data to distribute heavy keys.

## **5. Aggregate then join**

- a. Pre-aggregate big group values to reduce volume.



## **36. How to compute quantiles and median reliably at scale?**

### **Methods**

#### **1. approxQuantile (fast, approximate)**

```
quantiles = df.approxQuantile("salary", [0.25, 0.5,  
0.75], 0.01)
```

## 2. **percentile\_approx** (SQL function)

```
from pyspark.sql.functions import expr  
df.selectExpr("percentile_approx(salary, 0.5) as  
median").show()
```

## 3. **Exact median** (costly)

- Sort and compute middle row(s) — requires full shuffle and is expensive; typically not used on massive datasets.

## Tradeoffs

- Use approximate for speed (bounded error).
- Use exact only when dataset is small or exactness is mandatory.

## 37. Explain Bucketing: why, when and how (with code).

### Concept

Bucketing writes data into a fixed number of files (buckets) based on a hash of one or more columns. It improves join performance when both tables are bucketed on the same key and number of buckets match.

### How to write bucketed table

```
df.write.bucketBy(50,  
"id").sortBy("id") .mode ("over write") .saveAsTable("bu  
cketed_table")
```

### Benefits

- Avoids shuffle on joins if both sides bucketed identically.
- Faster joins for repeatable static datasets.

### Limitations

- Works best for static datasets (not for streaming).

- Need same number of buckets & same bucketing column for both tables.
- More complex to manage with schema evolution.

## 38. Explain Spark UI — key tabs and what to look for when tuning.

### Key UI areas

- **Jobs tab:** high-level job breakdown.
- **Stages tab:** identifies wide vs narrow stages; shows shuffle read/write.
- **Executors tab:** per-executor memory, GC times, task times.
- **SQL tab** (if enabled): physical plans, executed queries.
- **Storage tab:** cached RDD/DataFrame partitions.

### What to inspect

- Long-running tasks → skew.
- Large shuffle read/writes → reduce shuffles or optimize joins.
- Long GC times → tune executor memory or use MEMORY\_AND\_DISK\_SER.

- Skewed partition sizes → repartition or salt.
- Number of tasks vs cores → parallelism mismatch.

### Interview tip

Walk interviewer through an example: identify a slow job → check stages with the largest shuffle → inspect top tasks → hypothesize fixes (broadcast, repartition, reduce shuffle partitions).

## 39. What is Kryo serialization and why use it?

### Concept

Kryo is a faster, compact binary serializer alternative to Java serialization. It reduces serialization/deserialization overhead and network IO.

### How to enable

```
spark.conf.set("spark.serializer",
"org.apache.spark.serializer.KryoSerializer")
spark.conf.set("spark.kryo.registrationRequired",
```

```
"false" ) #optional
```

## When to use

- When you serialize large custom objects (case classes/pickles).
- For performance-sensitive workloads across the network.

## Caveats

- You may register custom classes for maximum speed and smaller serialized size.
- Debugging serialized objects is harder than Java serialization.

# ✓ 40. How to design idempotent pipelines and why it matters?

## Why idempotency

- Reprocessing shouldn't duplicate or corrupt downstream state.
- Essential for retries, failure recovery, and exactly-once semantics.

## Techniques

- Use **UPSERT (MERGE)** into Delta tables keyed by unique id.
- Use **transactional sinks** (Delta, ACID stores).
- Use **deduplication** via unique keys + latest timestamp.
- Use **atomic renames** when writing output files: write to temp dir → atomically move.

## Example (Delta upsert)

```
deltaTable.alias("t").merge(  
    updates.alias("u"),  
    "t.id = u.id"  
).whenMatchedUpdateAll().whenNotMatchedInsertAll().execute()
```

# ✓ 41. How to build ML feature pipelines in PySpark (high-level)?

## Components

- **Data ingestion** → reading structured + unstructured sources.

- **Cleaning & Imputation** → fillna, outlier handling.
- **Feature extraction** → StringIndexer, OneHotEncoder, Tokenizer, TF-IDF.
- **Feature scaling** → StandardScaler, MinMaxScaler.
- **Pipelines API** → combine transformers & estimators.

### Example

```
from pyspark.ml import Pipeline
from pyspark.ml.feature import StringIndexer,
OneHotEncoder, VectorAssembler, StandardScaler

si = StringIndexer(inputCol="cat",
outputCol="cat_idx")
ohe = OneHotEncoder(inputCols=["cat_idx"],
outputCols=["cat_vec"])
va =
VectorAssembler(inputCols=["num1","num2","cat_vec"],
outputCol="features")
scaler = StandardScaler(inputCol="features",
outputCol="scaled_features")
pipeline = Pipeline(stages=[si, ohe, va, scaler])
model = pipeline.fit(df)
```

```
df_transformed = model.transform(df)
```

### Tips

- Use VectorAssembler to prepare MLlib features.
- Persist intermediate large datasets only when reused.
- Use MLflow or Model Registry for artifacts.

## ✓ 42. How to do SCD Type 2 (Slowly

### Changing Dimensions) in PySpark?

#### Concept

Maintain historical records with `start_date`, `end_date`, and `is_current` flags when dimension attributes change.

#### Steps (high-level)

1. Read existing dimension Delta table.
2. Read incoming updates.
3. Identify rows to expire (existing rows where keys match but attributes differ).
4. Update existing rows `end_date` and `is_current=false`.

5. Insert new rows with `start_date=now`, `is_current=true`.

**Example (Delta merge pattern)**

```
from delta.tables import DeltaTable
from pyspark.sql.functions import current_timestamp

updates = new_data.withColumn("start_date",
    current_timestamp()).withColumn("is_current",
    lit(True))
dt = DeltaTable.forPath(spark, "/delta/dim_customer")

dt.alias("t").merge(
    updates.alias("u"),
    "t.customer_id = u.customer_id AND t.is_current =
true"
).whenMatchedUpdate(
    condition = "t.attr1 != u.attr1 OR t.attr2 !=
u.attr2",
    set = {"end_date": "current_timestamp()", "is_current": "false"}
).whenNotMatchedInsertAll().execute()
```

## Tips

- Maintain audit columns (etl\_ts, source, batch\_id).
- Compact table periodically to remove old files and optimize queries.

## 43. How to avoid the small files

### problem when writing Parquet/Delta?

#### Problem

Many small files cause high metadata overhead and slow reads.

#### Solutions

- **Write with coalesce/repartition:** reduce number of output files.

```
df.repartition(10).write.parquet("/path")
```

- **Use larger partition sizes:** avoid partitioning at too granular a level.
- **Use maxRecordsPerFile** when writing Parquet.

```
df.write.option("maxRecordsPerFile",  
1000000).parquet("/path")
```

- **Compaction (OPTIMIZE)**: Delta OPTIMIZE or Spark job to compact small files.
- **Streaming: use batch windows to write fewer larger files.**

## ✓ 44. How to perform schema drift

### handling in streaming pipelines?

#### Issues

Source schema can change (new fields, type changes) causing downstream failures.

#### Strategies

- **Use schema-on-read**: read JSON as string and parse with from\_json using evolving schema.
- **Schema registry**: enforce Avro/Schema registry for Kafka producers/consumers.
- **Schema evolution with Delta**: allow schema merge:

```
spark.conf.set("spark.databricks.ck.enabled", "true")
```

- **Versioned schema handling:** parse versions inside payload and transform accordingly.

## 45. Explain how to test PySpark jobs

### locally and in CI.

#### Local testing

- Use spark-submit with local master: --master local[\*].
- Use small sample datasets with representative edge cases.

#### Unit testing

- Use pytest + pyspark.sql.SparkSession fixture.
- Test transformations functionally: assert schemas, row counts, sample values.

#### CI best practices

- Use containerized Spark (Docker) or lightweight Spark images.

- Run tests on representative sample, not full dataset.
- Mock external dependencies (S3, Kafka) or use in-memory/local test doubles.

### Example pytest fixture

```
import pytest
from pyspark.sql import SparkSession

@pytest.fixture(scope="session")
def spark():
    spark =
        SparkSession.builder.master("local[2]").appName("test")
    yield spark
    spark.stop()
```

## 46. How to monitor and alert Spark jobs in production?

### Monitoring tools

- Spark UI (live)
- Ganglia / Prometheus / Grafana for metrics
- Cloud provider logs (Databricks, EMR, GCP Dataproc dashboards)
- Structured logging into ELK stack for driver/executor logs

### Key metrics to alert on

- Job duration > SLA
- Task failure rate > threshold
- GC time high or executor OOM
- Queue length in cluster manager
- Lag in streaming (processing time >> trigger interval)

### Practices

- Emit custom metrics (via Dropwizard/Prometheus) from executors.
- Centralize logs and set alerts on patterns (Exceptions, OOM).

- Use lineage & run metadata for observability (Airflow metadata, run IDs).



## 47. How to securely handle secrets, credentials in Spark jobs?

Do NOT hardcode secrets in code. Use:

- **Cluster secret stores** (Databricks Secrets, AWS Secrets Manager, Azure Key Vault).
- **Environment variables and instance roles** (IAM roles) for cloud storage access.
- **Encrypted config files** and read them at runtime from secure stores.
- **Short-lived tokens** for services.

### Example (AWS)

- Use **IAM role for EC2/EMR** so no AWS keys in code.
- Or fetch credentials at runtime from AWS Secrets Manager securely.

## 48. Explain how Spark integrates with Hive / external metastore.

### Integration points

- Spark can use Hive Metastore for table/catalog metadata.
- `spark.sql.catalogImplementation` can be set to `hive`.
- Use `enableHiveSupport()` in `SparkSession` to access Hive tables.

### Example

```
spark =  
SparkSession.builder.enableHiveSupport().getOrCreate()  
  
spark.sql("select * from db.table").show()
```

### Notes

- External metastore lets multiple engines (Presto, Hive, Spark) share metadata.
- Manage Hive-compatible partitioning and file formats (Parquet/ORC).



## 49. How to migrate PySpark jobs from on-prem to cloud (high-level checklist)?

### Checklist

1. **Inventory:** list jobs, dependencies, data sources, connectors.
2. **Data Storage:** move to S3/ADLS/GCS or connect via VPC.
3. **Secrets & IAM:** adopt cloud-native auth (IAM roles, key vaults).
4. **Cluster sizing:** choose instance types and autoscaling policies.
5. **Storage format:** use Parquet/Delta; consider converting CSV/JSON.
6. **Networking:** set up VPC/Subnets, private endpoints for S3.
7. **CI/CD:** deploy with IaC (Terraform), containerize where needed.
8. **Monitoring:** integrate cloud logging & metrics.
9. **Performance testing:** run scale tests on target cloud resources.
10. **Cost optimization:** spot instances/preemptible, right-sizing.

## **50. Top 10 Spark configs every data engineer should know (and why)**

1. `spark.sql.shuffle.partitions` – controls number of shuffle partitions (affects parallelism & small-files).
2. `spark.executor.memory` — memory per executor (avoid OOM or underutilization).
3. `spark.executor.cores` — cores per executor (controls task parallelism).
4. `spark.driver.memory` — driver memory for collect/joins on driver side.
5. `spark.serializer` — Kryo for faster serialization.
6. `spark.sql.adaptive.enabled` — enable AQE for runtime optimizations.
7. `spark.sql.autoBroadcastJoinThreshold` – broadcast join size threshold.
8. `spark.default.parallelism` – default number of partitions for RDD operations.
9. `spark.speculation` — speculative execution for slow tasks (defaults off, useful with skew).

10. spark.local.dir tmp dirs for shuffle and spill.



**Let's build your Data  
Engineering journey  
together!**



Call us directly at: 9989454737



<https://seekhobigdata.com/>

